

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatasience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatasience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454

Number of users: 256,059

Number of products: 74,258

Timespan: Oct 1999 - Oct 2012

Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review

10. Text - text of the review

Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative?

[Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

[1]. Reading Data

[1.1] Loading the data

The dataset is available in two forms

1. .csv file
2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: %matplotlib inline
import warnings
```

```
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

C:\ProgramData\Anaconda3\lib\site-packages\gensim\utils.py:1197: UserWarning: detected Windows; aliasing chunkize to chunkize_serial
  warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")
```

In [2]: *# using SQLite Table to read data.*

```

con = sqlite3.connect('C:/Users/Excel/Desktop/vins/database.sqlite')

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 50
0000 data points
# you can change the number to any other number based on your computing
power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Sco
re != 3 LIMIT 500000""", con)
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score
!= 3 LIMIT 50000""", con)

# Give reviews with Score>3 a positive rating(1), and reviews with a sc
ore<3 a negative rating(0).
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (50000, 10)

Out[2]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
--	----	-----------	--------	-------------	----------------------	------------------------

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1

```
In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)
```

```
In [4]: print(display.shape)
display.head()

(80668, 7)
```

Out[4]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...	2
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...	3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...	2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...	3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...	2

In [5]: display[display['UserId']=='AZY10LLTJ71NX']

Out[5]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT
--	--------	-----------	-------------	------	-------	------	-------

	UserId	ProductId	ProfileName	Time	Score	Text
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommended to try green tea extract to ...

In [6]: `display['COUNT(*)'].sum()`

Out[6]: 393063

[2] Exploratory Data Analysis

[2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

In [7]: `display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()`

Out[7]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfuln
--	----	-----------	--------	-------------	----------------------	----------

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfuln
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8)

ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True,
inplace=False, kind='quicksort', na_position='last')
```

```
In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time",
"Text"}, keep='first', inplace=False)
final.shape
```

```
Out[9]: (46072, 10)
```

```
In [10]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

```
Out[10]: 92.144
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [11]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)

display.head()
```

Out[11]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfuln
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	1
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	2

```
In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [13]: final["Time"] = pd.to_datetime(final["Time"], unit = "s")
final = final.sort_values(by = "Time")
```

```
In [14]: #Before starting the next phase of preprocessing lets see the number of
entries left
```

```
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?  
final['Score'].value_counts()
```

```
(46071, 10)
```

```
Out[14]: 1    38479  
        0     7592  
        Name: Score, dtype: int64
```

[3] Preprocessing

[3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [15]: # printing some random reviews  
sent_0 = final['Text'].values[0]  
print(sent_0)
```

```

print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)

```

This was a really good idea and the final product is outstanding. I use the decals on my car window and everybody asks where i bought the decals i made. Two thumbs up!

=====

Speaking as another Texan, I think the first rule for these delicious treats is to NOT order them during spring or summer. In fact, your safest bet is to ONLY order them in the dead of winter. LOL! As long as you do that, be prepared for a truly amazing treat! This package comes with 12 bite-sized delicacies. The chocolate is high-quality, the nuts are crunchy, and the overall taste couldn't be better. Definitely worth the price!

=====

Fast, easy and definitely delicious. Makes a great cup of coffee and very easy to make. Good purchase. Will continue to order from here.
Thanx...

=====

Naturally this review is based upon my cat's intake of Petite Cuisine. She's not a particular picky eater, so I can't say much about that. However, she looks to really enjoy this brand of cat food. I have tried some brands of wet food in the past that have made her sick (I know cats seem to have digestive systems that are prone to upsetting!) Petite Cuisine did not have any effect there, and she really enjoyed all the flavors. I don't feed her wet food often, usually just some tuna fish now and then. So, although this food is expensive if you used it at every meal, it is priced about the same as tuna, so it fits my needs perfectly. I'm sure my cat will enjoy the rest of this case, and I can keep the cat

nned tuna to myself for now :-)

=====

```
In [16]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

This was a really good idea and the final product is outstanding. I use the decals on my car window and everybody asks where i bought the decal s i made. Two thumbs up!

```
In [17]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all-tags-from-an-element
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)
```

This was a really good idea and the final product is outstanding. I use the decals on my car window and everybody asks where i bought the decals i made. Two thumbs up!

Speaking as another Texan, I think the first rule for these delicious treats is to NOT order them during spring or summer. In fact, your safest bet is to ONLY order them in the dead of winter. LOL! As long as you do that, be prepared for a truly amazing treat! This package comes with 12 bite-sized delicacies. The chocolate is high-quality, the nuts are crunchy, and the overall taste couldn't be better. Definitely worth the price!

Fast, easy and definitely delicious. Makes a great cup of coffee and very easy to make. Good purchase. Will continue to order from here. Thanks...

Naturally this review is based upon my cat's intake of Petite Cuisine. She's not a particular picky eater, so I can't say much about that. However, she looks to really enjoy this brand of cat food. I have tried some brands of wet food in the past that have made her sick (I know cats seem to have digestive systems that are prone to upsetting!) Petite Cuisine did not have any effect there, and she really enjoyed all the flavors. I don't feed her wet food often, usually just some tuna fish now and then. So, although this food is expensive if you used it at every meal, it is priced about the same as tuna, so it fits my needs perfectly. I'm sure my cat will enjoy the rest of this case, and I can keep the canned tuna to myself for now :-)

In [18]: `# https://stackoverflow.com/a/47091490/4084039`

```
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
```

```

phrase = re.sub(r"\ 's", " is", phrase)
phrase = re.sub(r"\ 'd", " would", phrase)
phrase = re.sub(r"\ 'll", " will", phrase)
phrase = re.sub(r"\ 't", " not", phrase)
phrase = re.sub(r"\ 've", " have", phrase)
phrase = re.sub(r"\ 'm", " am", phrase)
return phrase

```

```

In [19]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)

```

Fast, easy and definitely delicious. Makes a great cup of coffee and very easy to make. Good purchase. Will continue to order from here.
Thanx...

=====

```

In [20]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub(r"\S*\d\S*", "", sent_0).strip()
print(sent_0)

```

This was a really good idea and the final product is outstanding. I use the decals on my car window and everybody asks where i bought the decals i made. Two thumbs up!

```

In [21]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub(r'[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)

```

Fast easy and definitely delicious Makes a great cup of coffee and very easy to make Good purchase Will continue to order from here br Thanx

```

In [22]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have revmoved in

```


the 1st step

```
stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'o  
urs', 'ourselves', 'you', "you're", "you've",\  
"you'll", "you'd", 'your', 'yours', 'yourself', 'yoursele  
s', 'he', 'him', 'his', 'himself', \  
'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'it  
s', 'itself', 'they', 'them', 'their',\  
'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'th  
is', 'that', "that'll", 'these', 'those', \  
'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'h  
ave', 'has', 'had', 'having', 'do', 'does', \  
'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or',  
'because', 'as', 'until', 'while', 'of', \  
'at', 'by', 'for', 'with', 'about', 'against', 'between',  
'into', 'through', 'during', 'before', 'after',\  
'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out',  
'on', 'off', 'over', 'under', 'again', 'further',\  
'then', 'once', 'here', 'there', 'when', 'where', 'why', 'h  
ow', 'all', 'any', 'both', 'each', 'few', 'more',\  
'most', 'other', 'some', 'such', 'only', 'own', 'same', 's  
o', 'than', 'too', 'very', \  
's', 't', 'can', 'will', 'just', 'don', "don't", 'should',  
"should've", 'now', 'd', 'll', 'm', 'o', 're', \  
've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't",  
'didn', "didn't", 'doesn', "doesn't", 'hadn',\  
"hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "is  
n't", 'ma', 'mightn', "mightn't", 'mustn',\  
"mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn',  
"shouldn't", 'wasn', "wasn't", 'weren', "weren't", \  
'won', "won't", 'wouldn', "wouldn't"])
```

```
In [23]: # Combining all the above stundents  
from tqdm import tqdm  
preprocessed_reviews = []  
# tqdm is for printing the status bar  
for sentence in tqdm(final['Text'].values):  
    sentence = re.sub(r"http\S+", "", sentence)  
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
```

```

sentence = decontracted(sentence)
sentence = re.sub("\S*\d\S*", "", sentence).strip()
sentence = re.sub('[^A-Za-z]+', ' ', sentence)
# https://gist.github.com/sebleier/554280
sentence = ' '.join(e.lower() for e in sentence.split() if e.lower
() not in stopwords)
preprocessed_reviews.append(sentence.strip())

```

```

100%|████████████████████████████████████████| 46071/46071 [00:24<00:00, 184
5.69it/s]

```

In [27]: preprocessed_reviews[1500]

Out[27]: 'fast easy definitely delicious makes great cup coffee easy make good p
urchase continue order thanx'

[3.2] Preprocessing Review Summary

In [28]: *## Similarly you can do preprocessing for review summary also.*

[4] Featurization

[4.1] BAG OF WORDS

In [29]: X=preprocessed_reviews
Y=final["Score"]

In [30]: **from** sklearn.model_selection **import** train_test_split
#time based splitting
X_train,X_test,y_train,y_test = train_test_split(X,Y,test_size=0.33,shu
ffle=**False**)
X_train,X_cv,y_train,y_cv = train_test_split(X_train,y_train,test_size=
0.33,shuffle=**False**)

```
print(np.shape(X_train),y_train.shape)
print(np.shape(X_cv),y_cv.shape)
print(np.shape(X_test),y_test.shape)
```

```
(20680,) (20680,)
(10187,) (10187,)
(15204,) (15204,)
```

```
In [31]: from sklearn.feature_extraction.text import CountVectorizer
vectorizer=CountVectorizer(min_df=10,max_features=10000,ngram_range=(1,
2))
vectorizer.fit(X_train)
```

```
X_train_bow=vectorizer.transform(X_train)
X_cv_bow=vectorizer.transform(X_cv)
X_test_bow=vectorizer.transform(X_test)
```

```
print("After BOW vectorizer")
print(np.shape(X_train_bow),y_train.shape)
print(np.shape(X_cv_bow),y_cv.shape)
print(np.shape(X_test_bow),y_test.shape)
```

```
After BOW vectorizer
(20680, 10000) (20680,)
(10187, 10000) (10187,)
(15204, 10000) (15204,)
```

```
In [32]: from sklearn.model_selection import GridSearchCV,RandomizedSearchCV
from sklearn.tree import DecisionTreeClassifier
```

```
In [34]: depth=[1, 5, 10,20, 50, 100, 500, 1000]
min_split=[5,10,100,500]
parameter={'max_depth':depth,'min_samples_split':min_split}

dt=GridSearchCV(DecisionTreeClassifier(class_weight='balanced'),paramet
er,verbose=1,scoring='roc_auc')
dt.fit(X_train_bow,y_train)

opt_depth,opt_min_split=dt.best_params_.get('max_depth'),dt.best_params
```

```

_.get('min_samples_split')
print("best optimized depth:",opt_depth)
print("best optimized min_split:",opt_min_split)

train_score=dt.cv_results_.get("mean_train_score")
test_score=dt.cv_results_.get("mean_test_score")

plt.plot(np.arange(len(depth)),train_score[::4], 'r',label="train_data
(5)")
plt.plot(np.arange(len(depth)),test_score[::4], 'r--',label="cv_data(5)"
)

plt.plot(np.arange(len(depth)),train_score[1::4], 'g',label="train_data
(10)")
plt.plot(np.arange(len(depth)),test_score[1::4], 'g--',label="cv_data(1
0)")

plt.plot(np.arange(len(depth)),train_score[2::4], 'b',label="train_data
(100)")
plt.plot(np.arange(len(depth)),test_score[2::4], 'b--',label="cv_data(10
0)")

plt.plot(np.arange(len(depth)),train_score[3::4], 'y',label="train_data
(500)")
plt.plot(np.arange(len(depth)),test_score[3::4], 'y--',label="cv_data(50
0)")

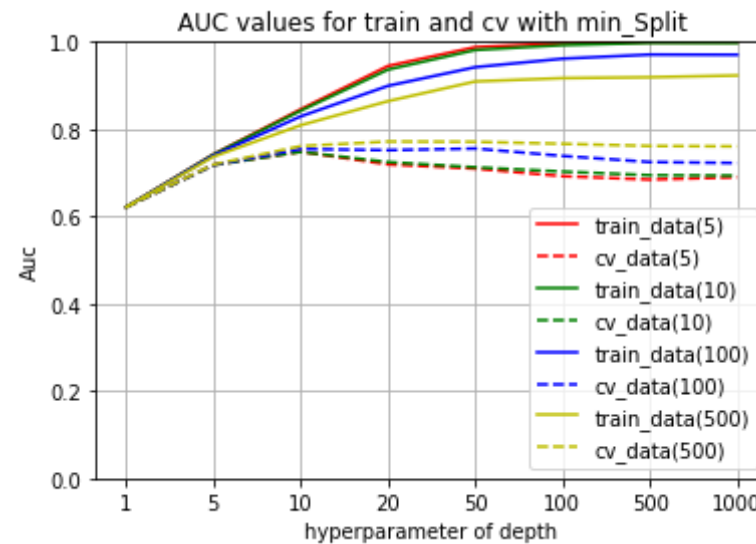
plt.xticks(np.arange(len(depth)),depth)
plt.legend()
plt.xlabel("hyperparameter of depth")
plt.ylabel("Auc")
plt.title("AUC values for train and cv with min_Split")
plt.ylim(0,1)
plt.grid(True)
plt.show()

```

Fitting 3 folds for each of 32 candidates, totalling 96 fits

[Parallel(n_jobs=1)]: Done 96 out of 96 | elapsed: 3.1min finished

best optimized depth: 20
best optimized min_split: 500





1. we get correct visualization of optimized max_depth.
2. we will go for the seaborn heatmap for min_samples_split.

```
In [72]: df_heatmap = pd. DataFrame(train_score. reshape(8, 4), index=depth, col
        umns=min_split )
        fig = plt. figure(figsize=(6, 3))
        heatmap = sns. heatmap(df_heatmap, annot=True)
        plt. ylabel('Depth' , size=18)
        plt. xlabel('min_Sample_Split' , size=18)
        plt. title("Train Data", size=24)
        plt. show

        df_heatmap = pd. DataFrame(test_score. reshape(8, 4), index=depth, colu
        mns=min_split )
        fig = plt. figure(figsize=(5, 3))
        heatmap = sns. heatmap(df_heatmap, annot=True)
        plt. ylabel('Depth' , size=18)
        plt. xlabel('min_Sample_Split' , size=18)
```

```
plt. title("Train Data", size=24)  
plt. show
```

Out[72]: <function matplotlib.pyplot.show(*args, **kw)>



```
In [53]: from sklearn.metrics import roc_auc_score  
dt=DecisionTreeClassifier(class_weight='balanced',max_depth=20,min_samp  
les_split=500)  
dt.fit(X_train_bow,y_train)
```

```


y_test_proba=dt.predict_proba(X_test_bow)
y_train_proba=dt.predict_proba(X_train_bow)

fpr, tpr, threshold=roc_curve(y_train, y_train_proba[:,1])
fpr1, tpr1, threshold1=roc_curve(y_test, y_test_proba[:,1])
print("AUC value for test data :", roc_auc_score(y_test, y_test_proba[:,1]))

plt.plot(fpr, tpr, 'r', label='train_data')
plt.plot(fpr1, tpr1, 'b', label='test data')
plt.legend(loc='upper right')
plt.grid(True)
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.title('AUC curve for train and test data')
plt.show()

```


AUC value for test data : 0.8110928038929697



Confusion matrix for BOW

```
In [54]: def confusion_matrix(y_test, test):
    from sklearn.metrics import confusion_matrix
    y_test_predict=dt.predict(test)
    cm_test=confusion_matrix(y_test,y_test_predict)
    import seaborn as sns
    class_label=['0','1']
    df_cm=pd.DataFrame(cm_test,index=class_label,columns=class_label)
    sns.heatmap(df_cm,annot=True,fmt='d')
    plt.xlabel("predicted label")
    plt.ylabel("actual label")
    plt.show()
```

```
In [55]: confusion_matrix(y_test,X_test_bow)
```



Top 20 importance features for positive and negative class for BOW

```
In [73]: # Please write all the code with proper documentation
dt = DecisionTreeClassifier(class_weight= 'balanced',max_depth=20, min_
samples_split=500)
```



```

dt.fit(X_train_bow,y_train)
feat_log = dt.feature_importances_

count_vect = CountVectorizer(min_df=10,max_features=10000,ngram_range=(
1,2))
s = count_vect.fit(X_train)
s = pd.DataFrame(feat_log.T,columns=['+ve'])
s['feature'] = count_vect.get_feature_names()

v = s.sort_values(by = '+ve',kind = 'quicksort',ascending= False)
print("Top 20  important features of positive class", np.array(v['feature'][:20]))
print("*****400")
print("Top 20  important features of negative class",np.array(v.tail(20)
['feature']))

```

Top 20 important features of positive class ['not' 'great' 'best' 'delicious' 'love' 'perfect' 'good' 'bad' 'loves' 'excellent' 'disappointed' 'nice' 'favorite' 'wonderful' 'thought' 'not good' 'worst' 'unfortunately' 'pleased' 'reviews']

Top 20 important features of negative class ['function' 'fully' 'fruits' 'fuller' 'fruits vegetables' 'fruits veggies' 'fruity' 'fruity taste' 'frustrated' 'frustrating' 'fry' 'frying' 'fudge' 'fuel' 'full' 'full bodied' 'full cup' 'full flavor' 'full flavored' 'zuke']

Graphviz for BOW

In [80]: `# https://github.com/scikit-learn/scikit-learn/issues/9952
 from IPython.display import Image`

```

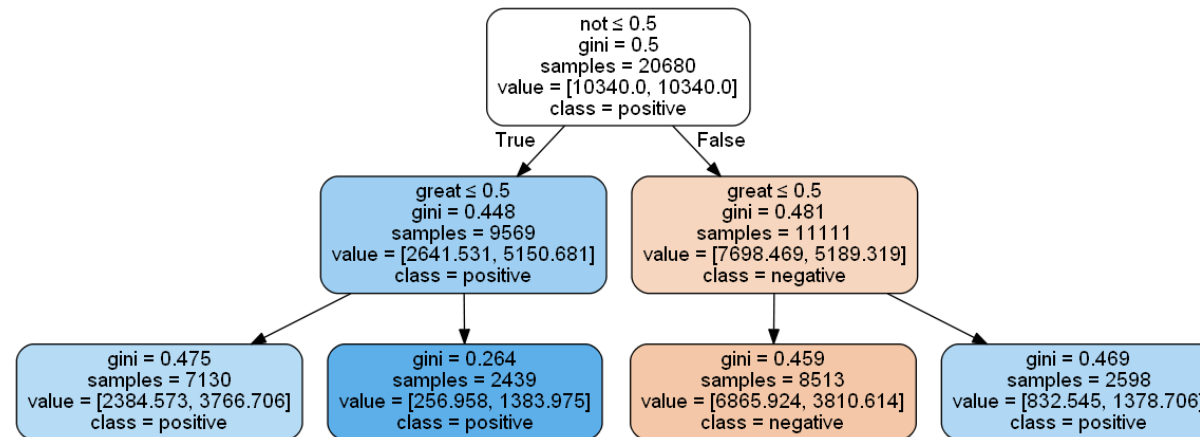
from sklearn.externals.six import StringIO
from sklearn.tree import export_graphviz
from sklearn.tree import DecisionTreeClassifier
import pydot
vectorizer=CountVectorizer(ngram_range=(1,2),min_df=10,max_features=100
00)
vectorizer.fit(X_train)
feature_name=vectorizer.get_feature_names()

dt=DecisionTreeClassifier(class_weight='balanced',max_depth=2,min_sampl
es_split=500)
dt.fit(X_train_bow,y_train)

dot_data=StringIO()
export_graphviz(dt,out_file=dot_data,feature_names=feature_name,class_n
ames=["negative",'positive'],filled=True,rounded=True,special_characters=True)
graph = pydot.graph_from_dot_data(dot_data.getvalue())[0]
Image(graph.create_png())

```

Out[80]:



[2] TF-IDF

```

In [140]: X=preprocessed_reviews
          Y=final["Score"]

```

```
In [141]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,Y,test_size=0.33,shuffle=False)
X_train,X_cv,y_train,y_cv=train_test_split(X_train,y_train,test_size=0.33,shuffle=False)
print(np.shape(X_train),y_train.shape)
print(np.shape(X_cv),y_cv.shape)
print(np.shape(X_test),y_test.shape)

(20680,) (20680,)
(10187,) (10187,)
(15204,) (15204,)
```

```
In [142]: from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer_tf=TfidfVectorizer(min_df=10,max_features=10000,ngram_range=(1,2))
vectorizer_tf.fit(X_train)

X_train_tf=vectorizer_tf.transform(X_train)
X_cv_tf=vectorizer_tf.transform(X_cv)
X_test_tf=vectorizer_tf.transform(X_test)

print("After TF-IDF vectorizer")
print(np.shape(X_train_tf),y_train.shape)
print(np.shape(X_cv_tf),y_cv.shape)
print(np.shape(X_test_tf),y_test.shape)

After TF-IDF vectorizer
(20680, 10000) (20680,)
(10187, 10000) (10187,)
(15204, 10000) (15204,)
```

```
In [96]: depth=[1,5,10,20,50,100,500,1000]
min_split=[5,10,100,500]
parameter={'max_depth':depth,"min_samples_split":min_split}
dt=GridSearchCV(DecisionTreeClassifier(class_weight='balanced'),parameter,verbose=1,scoring='roc_auc')
dt.fit(X_train_tf,y_train)
```

```

opt_depth,opt_min_sample_split=dt.best_params_.get("max_depth"),dt.best
_params_.get('min_samples_split')
print("The best optimized depth",opt_depth)
print("the best optimized depth",opt_min_sample_split)

train_score=dt.cv_results_.get("mean_train_score")
cv_score=dt.cv_results_.get('mean_test_score')

plt.plot(np.arange(len(depth)),train_score[::4],'r',label="train_data
(5)")
plt.plot(np.arange(len(depth)),test_score[::4],'r--',label="cv_data(5)"
)

plt.plot(np.arange(len(depth)),train_score[1::4],'g',label="train_data
(10)")
plt.plot(np.arange(len(depth)),test_score[1::4],'g--',label="cv_data(1
0)")

plt.plot(np.arange(len(depth)),train_score[2::4],'b',label="train_data
(100)")
plt.plot(np.arange(len(depth)),test_score[2::4],'b--',label="cv_data(10
0)")

plt.plot(np.arange(len(depth)),train_score[3::4],'y',label="train_data
(500)")
plt.plot(np.arange(len(depth)),test_score[3::4],'y--',label="cv_data(50
0)")

plt.xticks(np.arange(len(depth)),depth)
plt.legend()
plt.xlabel("hyperparameter of depth")
plt.ylabel("Auc")
plt.title("AUC values for train and cv with min_Split")
plt.ylim(0,1)
plt.grid(True)
plt.show

```


Fitting 3 folds for each of 32 candidates, totalling 96 fits

[Parallel(n_jobs=1)]: Done 96 out of 96 | elapsed: 4.6min finished

```
Parallel(n_jobs=1): Done 50 out of 50 | elapsed: 4.0min finished
```

The best optimized depth 20
the best optimized depth 500

Out[96]: <function matplotlib.pyplot.show(*args, **kw)>


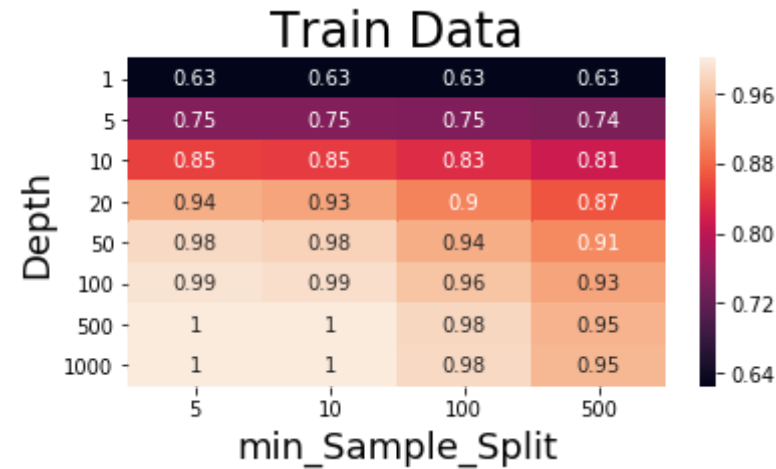


```
In [98]: df_heatmap = pd. DataFrame(train_score. reshape(8, 4), index=depth, col
umns=min_split )
fig = plt. figure(figsize=(6, 3))
heatmap = sns. heatmap(df_heatmap, annot=True)
plt. ylabel('Depth' , size=18)
plt. xlabel('min_Sample_Split' , size=18)
plt. title("Train Data", size=24)
plt. show

df_heatmap = pd. DataFrame(test_score. reshape(8, 4), index=depth, colu
mns=min_split )
fig = plt. figure(figsize=(6, 3))
heatmap = sns. heatmap(df_heatmap, annot=True)
plt. ylabel('Depth' , size=18)
plt. xlabel('min_Sample_Split' , size=18)
```

```
plt. title("Train Data", size=24)
plt. show
```

Out[98]: <function matplotlib.pyplot.show(*args, **kw)>



```
In [143]: from sklearn.metrics import roc_auc_score
dt=DecisionTreeClassifier(class_weight='balanced',max_depth=20,min_samples_split=500)
dt.fit(X_train_tf,y_train)
```

```


y_test_proba=dt.predict_proba(X_test_tf)
y_train_proba=dt.predict_proba(X_train_tf)

fpr, tpr, threshold=roc_curve(y_train, y_train_proba[:,1])
fpr1, tpr1, threshold1=roc_curve(y_test, y_test_proba[:,1])
print("AUC value for test data :", roc_auc_score(y_test, y_test_proba[:,1]))

plt.plot(fpr, tpr, 'r', label='train_data')
plt.plot(fpr1, tpr1, 'b', label='test data')
plt.legend(loc='upper right')
plt.grid(True)
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.title('AUC curve for train and test data')
plt.show()


```

AUC value for test data : 0.7985082357816222



Confusion matrix

```
In [99]: confusion_matrix(y_test,X_test_tf)
```



Top 20 features importance of positive and negative

```
In [104]: dt=DecisionTreeClassifier(class_weight='balanced',max_depth=20,min_samples_split=500)
dt.fit(X_train_tf,y_train)
feat_log=dt.feature_importances_

vectorizer=TfidfVectorizer(min_df=10,max_features=10000,ngram_range=(1,2))
s=vectorizer.fit_transform(X_train)
s=pd.DataFrame(feat_log.T,columns=["+ve"])
s['feature']=vectorizer.get_feature_names()
v=s.sort_values(by='+ve',kind='quicksort',ascending=False)
print('top 20 important feature of positive',np.array(v['feature'][:20]))
print("*"*300)
```



```
print('top 20 important feature of negative',np.array(v.tail(20)['feature']))
```

```
top 20 important feature of positive ['not' 'great' 'delicious' 'best'
'love' 'good' 'perfect' 'bad' 'loves'
'excellent' 'nice' 'wonderful' 'favorite' 'disappointed' 'tasty'
'thought' 'not good' 'reviews' 'easy' 'terrible']
*****
*****
*****
*****
*****
top 20 important feature of negative ['furthermore' 'furniture' 'frustrating'
'fur' 'fry' 'frying' 'fudge'
'fuel' 'full' 'full bodied' 'full cup' 'full flavor' 'full flavored'
'fuller' 'fully' 'fun' 'function' 'funky' 'funny' 'zuke']
```

Graphviz for TF-IDF

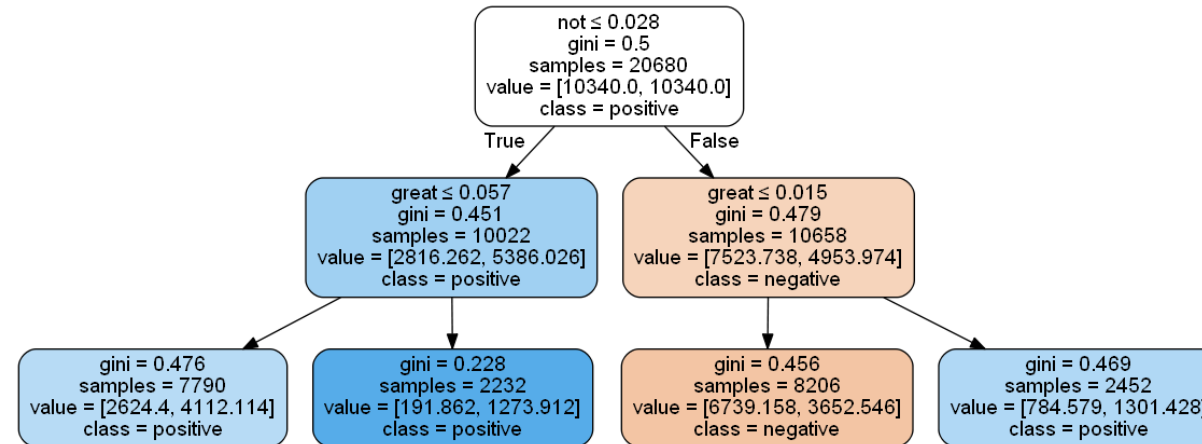
```
In [105]: # https://github.com/scikit-learn/scikit-learn/issues/9952
from IPython.display import Image
from sklearn.externals.six import StringIO
from sklearn.tree import export_graphviz
from sklearn.tree import DecisionTreeClassifier
import pydot
vectorizer=TfidfVectorizer(ngram_range=(1,2),min_df=10,max_features=1000)
vectorizer.fit(X_train)
feature_name=vectorizer.get_feature_names()

dt=DecisionTreeClassifier(class_weight='balanced',max_depth=2,min_samples_split=500)
dt.fit(X_train_tf,y_train)

dot_data=StringIO()
export_graphviz(dt,out_file=dot_data,feature_names=feature_name,class_names=["negative", 'positive'],filled=True,rounded=True,special_character
```

```
s=True)
graph = pydot.graph_from_dot_data(dot_data.getvalue())[0]
Image(graph.create_png())
```

Out[105]:



[3] AVG-W2V

```
In [131]: X=preprocessed_reviews
          Y=final['Score']
```

```
In [132]: from sklearn.model_selection import train_test_split
          X_train,X_test,y_train,y_test=train_test_split(X,Y,test_size=0.33,shuffle=False)
          X_train,X_cv,y_train,y_cv=train_test_split(X_train,y_train,test_size=0.33,shuffle=False)

          print(np.shape(X_train),y_train.shape)
          print(np.shape(X_cv),y_cv.shape)
          print(np.shape(X_test),y_test.shape)

          (20680,) (20680,)
          (10187,) (10187,)
          (15204,) (15204,)
```

```
In [133]: i=0
list_of_sentence=[]
for sentence in preprocessed_reviews:
    list_of_sentence.append(sentence.split())
```

```
In [134]: sent_of_train=[]
for sent in X_train:
    sent_of_train.append(sent.split())
```

```
In [135]: sent_of_cv=[]
for sent in X_cv:
    sent_of_cv.append(sent.split())

sent_of_test=[]
for sent in X_test:
    sent_of_test.append(sent.split())

# Train your own Word2Vec model using your own train text corpus
# min_count = 5 considers only words that occurred at least 5 times
w2v_model=Word2Vec(sent_of_train,min_count=5,size=50, workers=4)

w2v_words = list(w2v_model.wv.vocab)
```

```
In [136]: train_vectors = [];
for sent in sent_of_train:
    sent_vec = np.zeros(50)
    cnt_words =0;
    for word in sent: #
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    train_vectors.append(sent_vec)

cv_vectors = [];
for sent in sent_of_cv:
```

```

sent_vec = np.zeros(50)
cnt_words = 0;
for word in sent: #
    if word in w2v_words:
        vec = w2v_model.wv[word]
        sent_vec += vec
        cnt_words += 1
if cnt_words != 0:
    sent_vec /= cnt_words
cv_vectors.append(sent_vec)

# compute average word2vec for each review for X_test .
test_vectors = [];
for sent in sent_of_test:
    sent_vec = np.zeros(50)
    cnt_words = 0;
    for word in sent: #
        if word in w2v_words:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
    if cnt_words != 0:
        sent_vec /= cnt_words
    test_vectors.append(sent_vec)

```

```

In [137]: X_train_wv=train_vectors
          X_cv_wv=cv_vectors
          X_test_wv=test_vectors

```

```

In [116]: depth=[1,5,10,20,50,100,500,1000]
          min_split=[5,10,100,500]
          parameter={'max_depth':depth,"min_samples_split":min_split}
          dt=GridSearchCV(DecisionTreeClassifier(class_weight='balanced'),parameter,verbose=1,scoring='roc_auc')
          dt.fit(X_train_wv,y_train)

```

```

opt_depth,opt_min_sample_split=dt.best_params_.get("max_depth"),dt.best
_params_.get('min_samples_split')
print("The best optimized depth",opt_depth)
print("the best optimized depth",opt_min_sample_split)

train_score=dt.cv_results_.get("mean_train_score")
cv_score=dt.cv_results_.get('mean_test_score')

plt.plot(np.arange(len(depth)),train_score[::4],'r',label="train_data
(5)")
plt.plot(np.arange(len(depth)),test_score[::4],'r--',label="cv_data(5)"
)

plt.plot(np.arange(len(depth)),train_score[1::4],'g',label="train_data
(10)")
plt.plot(np.arange(len(depth)),test_score[1::4],'g--',label="cv_data(1
0)")

plt.plot(np.arange(len(depth)),train_score[2::4],'b',label="train_data
(100)")
plt.plot(np.arange(len(depth)),test_score[2::4],'b--',label="cv_data(10
0)")

plt.plot(np.arange(len(depth)),train_score[3::4],'y',label="train_data
(500)")
plt.plot(np.arange(len(depth)),test_score[3::4],'y--',label="cv_data(50
0)")

plt.xticks(np.arange(len(depth)),depth)
plt.legend()
plt.xlabel("hyperparameter of depth")
plt.ylabel("Auc")
plt.title("AUC values for train and cv with min_Split")
plt.ylim(0,1)
plt.grid(True)
plt.show

```


Fitting 3 folds for each of 32 candidates, totalling 96 fits

[Parallel(n_jobs=1)]: Done 96 out of 96 | elapsed: 1.5min finished

```
Parallel(n_jobs=1): Done 30 out of 30 | elapsed: 17.5min finished
```

The best optimized depth 10
the best optimized depth 500

Out[116]: <function matplotlib.pyplot.show(*args, **kw)>





```
In [117]: df_heatmap = pd. DataFrame(train_score. reshape(8, 4), index=depth, col
          umns=min_split )
          fig = plt. figure(figsize=(6, 3))
          heatmap = sns. heatmap(df_heatmap, annot=True)
          plt. ylabel('Depth' , size=18)
          plt. xlabel('min_Sample_Split' , size=18)
          plt. title("Train Data", size=24)
          plt. show

          df_heatmap = pd. DataFrame(test_score. reshape(8, 4), index=depth, colu
          mns=min_split )
          fig = plt. figure(figsize=(6, 3))
          heatmap = sns. heatmap(df_heatmap, annot=True)
          plt. ylabel('Depth' , size=18)
          plt. xlabel('min_Sample_Split' , size=18)
```

```
plt. title("Train Data", size=24)
plt. show
```

Out[117]: <function matplotlib.pyplot.show(*args, **kw)>



```
In [138]: from sklearn.metrics import roc_auc_score
dt=DecisionTreeClassifier(class_weight='balanced',max_depth=10,min_samp
les_split=500)
dt.fit(X_train_wv,y_train)
```

```


y_test_proba=dt.predict_proba(X_test_wv)
y_train_proba=dt.predict_proba(X_train_wv)

fpr, tpr, threshold=roc_curve(y_train, y_train_proba[:,1])
fpr1, tpr1, threshold1=roc_curve(y_test, y_test_proba[:,1])
print("AUC value for test data :", roc_auc_score(y_test, y_test_proba[:,1]))

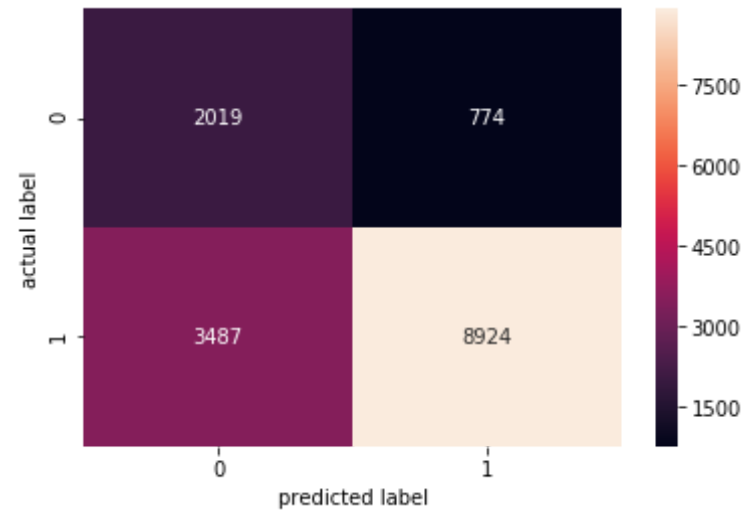
plt.plot(fpr, tpr, 'r', label='train_data')
plt.plot(fpr1, tpr1, 'b', label='test data')
plt.legend(loc='upper right')
plt.grid(True)
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.title('AUC curve for train and test data')
plt.show()

```

AUC value for test data : 0.7936668766544397



In [139]: `confusion_matrix(y_test, X_test_wv)`



TF-IDF W2V

```
In [119]: X=preprocessed_reviews  
Y=final["Score"]
```

```
In [120]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33, shuffle=False) # this is random splitting  
X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.33, shuffle=False)
```

```
In [121]: model = TfidfVectorizer()  
tfidf_matrix = model.fit_transform(X_train)  
# we are converting a dictionary with word as a key, and the idf as a value  
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

```
In [122]: tfidf_feat = model.get_feature_names() # tfidf words/col-names  
# final_tf_idf is the sparse matrix with row= sentence, col=word and ce
```

```

ll_val = tfidf

tfidf_train_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(sent_of_train): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_train_vectors.append(sent_vec)
    row += 1

```

```

100%|████████████████████████████████████████| 20680/20680 [05:03<00:00, 68.09it/s]

```

```

In [123]: tfidf_cv_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(sent_of_cv): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are

```

```

        # dictionary[word] = idf value of word in whole corpus
        # sent.count(word) = tf value of word in this review
        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vec * tf_idf)
        weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_cv_vectors.append(sent_vec)
    row += 1

```

100%|██| 10187/10187 [02:28<00:00, 68.53it/s]

In [124]:

```

tfidf_test_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(sent_of_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            # tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_test_vectors.append(sent_vec)
    row += 1

```

100%|██| 15204/15204 [03:23<00:00, 74.55it/s]

In [125]: X_train_tw=tfidf_train_vectors

```
X_cv_tw=tfidf_cv_vectors
X_test_tw=tfidf_test_vectors
```

```
In [126]: depth=[1,5,10,20,50,100,500,1000]
min_split=[5,10,100,500]
parameter={'max_depth':depth,"min_samples_split":min_split}
dt=GridSearchCV(DecisionTreeClassifier(class_weight='balanced'),parameter,verbose=1,scoring='roc_auc')
dt.fit(X_train_tw,y_train)

opt_depth,opt_min_sample_split=dt.best_params_.get("max_depth"),dt.best_params_.get('min_samples_split')
print("The best optimized depth",opt_depth)
print("the best optimized depth",opt_min_sample_split)

train_score=dt.cv_results_.get("mean_train_score")
cv_score=dt.cv_results_.get('mean_test_score')

plt.plot(np.arange(len(depth)),train_score[::4],'r',label="train_data(5)")
plt.plot(np.arange(len(depth)),test_score[::4],'r--',label="cv_data(5)")

plt.plot(np.arange(len(depth)),train_score[1::4],'g',label="train_data(10)")
plt.plot(np.arange(len(depth)),test_score[1::4],'g--',label="cv_data(10)")

plt.plot(np.arange(len(depth)),train_score[2::4],'b',label="train_data(100)")
plt.plot(np.arange(len(depth)),test_score[2::4],'b--',label="cv_data(100)")

plt.plot(np.arange(len(depth)),train_score[3::4],'y',label="train_data(500)")
plt.plot(np.arange(len(depth)),test_score[3::4],'y--',label="cv_data(500)")

plt.xticks(np.arange(len(depth)),depth)
```


```
plt.legend()
plt.xlabel("hyperparameter of depth")
plt.ylabel("Auc")
plt.title("AUC values for train and cv with min_Split")
plt.ylim(0,1)
plt.grid(True)
plt.show
```

Fitting 3 folds for each of 32 candidates, totalling 96 fits

[Parallel(n_jobs=1)]: Done 96 out of 96 | elapsed: 1.4min finished

The best optimized depth 20
the best optimized depth 500

Out[126]: <function matplotlib.pyplot.show(*args, **kw)>





```
In [127]: df_heatmap = pd. DataFrame(train_score. reshape(8, 4), index=depth, col
umns=min_split )
fig = plt. figure(figsize=(6, 3))
heatmap = sns. heatmap(df_heatmap, annot=True)
plt. ylabel('Depth' , size=18)
plt. xlabel('min_Sample_Split' , size=18)
```

```
plt. title("Train Data", size=24)
plt. show

df_heatmap = pd. DataFrame(test_score. reshape(8, 4), index=depth, columns=min_split )
fig = plt. figure(figsize=(6, 3))
heatmap = sns. heatmap(df_heatmap, annot=True)
plt. ylabel('Depth' , size=18)
plt. xlabel('min_Sample_Split' , size=18)
plt. title("Train Data", size=24)
plt. show
```

Out[127]: <function matplotlib.pyplot.show(*args, **kw)>






```
In [130]: from sklearn.metrics import roc_auc_score
dt=DecisionTreeClassifier(class_weight='balanced',max_depth=20,min_samples_split=500)
dt.fit(X_train_tw,y_train)

y_test_proba=dt.predict_proba(X_test_tw)
y_train_proba=dt.predict_proba(X_train_tw)


fpr,tpr,threshold=roc_curve(y_train,y_train_proba[:,1])
fpr1,tpr1,threshold1=roc_curve(y_test,y_test_proba[:,1])
print("AUC value for test data :",roc_auc_score(y_test,y_test_proba[:,1]))

plt.plot(fpr,tpr,'r',label='train_data')
plt.plot(fpr1,tpr1,'b',label='test data')
plt.legend(loc='upper right')
plt.grid(True)
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.title('AUC curve for train and test data')
plt.show()
```

AUC value for test data : 0.7494858703672981



```
In [129]: confusion_matrix(y_test,X_test_tw)
```



Conclusion


```
In [149]: from tabulate import tabulate
print(tabulate ([['BOW(20)', 500, 81],['TF-IDF(20)',500,79],['AVG-W2V(10)',500,79] , ['TFIDF-W2V(20)',500,74]], headers=['Vectorizer(opt_depth)', 'best_min_sample_split','AUC_test']))
```

Vectorizer(opt_depth)	best_min_sample_split	AUC_test
BOW(20)	500	81
TF-IDF(20)	500	79
AVG-W2V(10)	500	79
TFIDF-W2V(20)	500	74

1. cost of computation is fast.
2. The best model is Bag of words.
3. we can improve the model by taking more datapoints.

[5] Assignment 8: Decision Trees

1. Apply Decision Trees on these feature sets

- **SET 1:**Review text, preprocessed one converted into vectors using (BOW)
- **SET 2:**Review text, preprocessed one converted into vectors using (TFIDF)
- **SET 3:**Review text, preprocessed one converted into vectors using (AVG W2v)
- **SET 4:**Review text, preprocessed one converted into vectors using (TFIDF W2v)

2. The hyper paramter tuning (best `depth` in range [1, 5, 10, 50, 100, 500, 100], and the best `min_samples_split` in range [5, 10, 100, 500])

- Find the best hyper parameter which will give the maximum [AUC](#) value
- Find the best hyper paramter using k-fold cross validation or simple cross validation data
- Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task of hyperparameter tuning

3. Graphviz

- Visualize your decision tree with Graphviz. It helps you to understand how a decision is being made, given a new vector.
- Since feature names are not obtained from word2vec related models, visualize only BOW & TFIDF decision trees using Graphviz
- Make sure to print the words in each node of the decision tree instead of printing its index.
- Just for visualization purpose, limit max_depth to 2 or 3 and either embed the generated images of graphviz in your notebook, or directly upload them as .png files.



4. Feature importance

- Find the top 20 important features from both feature sets **Set 1** and **Set 2** using `feature_importances_` method of [Decision Tree Classifier](#) and print their corresponding feature names

5. Feature engineering

- To increase the performance of your model, you can also experiment with with feature engineering like :
 - Taking length of reviews as another feature.
 - Considering some features from review summary as well.

6. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure.
 -  Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.
 -  Along with plotting ROC curve, you need to print the [confusion matrix](#) with predicted and original labels of test data points. Please

visualize your confusion matrices using [seaborn heatmaps](#).



7. [Conclusion](#)

- [You need to summarize the results at the end of the notebook, summarize it in the table format. To print out a table please refer to this prettytable library link](#)



Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method `fit_transform()` on you train data, and apply the method `transform()` on cv/test data.
4. For more details please go through this [link](#).

Applying Decision Trees

[5.1] Applying Decision Trees on BOW, SET 1

```
In [0]: # Please write all the code with proper documentation
```

[5.1.1] Top 20 important features from SET 1

```
In [0]: # Please write all the code with proper documentation
```

[5.1.2] Graphviz visualization of Decision Tree on BOW, SET 1

In [0]: *# Please write all the code with proper documentation*

[5.2] Applying Decision Trees on TFIDF, SET 2

In [0]: *# Please write all the code with proper documentation*

[5.2.1] Top 20 important features from SET 2

In [0]: *# Please write all the code with proper documentation*

[5.2.2] Graphviz visualization of Decision Tree on TFIDF, SET 2

In [0]: *# Please write all the code with proper documentation*

[5.3] Applying Decision Trees on AVG W2V, SET 3

In [0]: *# Please write all the code with proper documentation*

[5.4] Applying Decision Trees on TFIDF W2V, SET 4

In [0]: *# Please write all the code with proper documentation*

[6] Conclusions

In [0]: *# Please compare all your models using Prettytable library*